

From Signals to Sovereignty: A Technical Guide to Host-Local Biophysical Authorship

Designing the HostInterfaceBus: A Standardized Signal Schema for Biological Inputs

The foundational step toward achieving biophysical autonomy is the creation of a standardized communication layer that normalizes the diverse array of signals generated by the human body. This layer, which we will call the **HostInterfaceBus**, serves as the universal protocol through which all assistive devices, whether current or future, must communicate. Its purpose is to abstract away the idiosyncrasies of individual sensors and signal types, presenting a consistent, typed, and event-driven stream of information to the host-node's processing logic . By defining a clear schema, this bus transforms the user's biological system into a stable reference point, where new hardware simply becomes another producer of well-defined events rather than a unique challenge requiring bespoke integration . The design of this bus prioritizes interoperability, precision, and extensibility, ensuring that the user's biology remains the sovereign operating system of their own digital life.

The specification for the **HostInterfaceBus** must be designed to handle multiple input modalities, ranging from primary, accessible sources like surface electromyography (sEMG), motion sensors, and voice features, to optional, higher-bandwidth inputs such as electroencephalography (EEG) and eye-tracking . Each modality provides a different window into the user's state and intent, and the bus must be capable of ingesting, timestamping, and labeling data from all of them. To achieve this, the bus should define a canonical event structure, which in Rust could be represented as a **BciEvent** struct. This struct would encapsulate all necessary metadata to make the signal unambiguous and actionable. Key fields for this struct include a high-resolution **timestamp_ns** field, typically using a **u64** integer, to enable precise synchronization across different sensor streams; a **source_id** string to uniquely identify the physical device or sensor generating the data; a **channel_id** string to specify the particular type of measurement being reported; a normalized **value** field to hold the quantitative result of the measurement; and a **confidence** score, a **f32** floating-point number between 0.0 and

1.0, to indicate the decoder's certainty in the reading, which is crucial for filtering out noisy or uncertain data [26](#).

The `channel_id` is particularly critical as it standardizes the feature identifiers themselves. Drawing from extensive research into psychophysiological metrics, a comprehensive list of potential channels can be defined, covering cerebral, ocular, cardiovascular, dermal, and other task-performance related features [19](#). For instance, under the 'Cerebral' category, one might find channels like `EEG_ALPHA_POWER` or `HEMODYNAMIC_RESPONSE_FNIRS`. Under 'Ocular', channels could include `PUPIL_DIAMETER` or `FIXATION_COUNT`, while 'Cardiovascular' might have `HEART_RATE` or `RESPIRATION_RATE` [19](#). Creating a master registry of these standardized channel IDs ensures that any device manufacturer or software developer can map their proprietary data to a common vocabulary, facilitating seamless integration. This approach aligns with frameworks that seek to operationalize neuro-symbolic architectures by establishing clear protocols for data exchange [2](#).

Once low-level events are flowing through the bus, the next function is to aggregate and interpret them into high-level semantic intentions. This process decouples the physical act (e.g., a muscle contraction) from the logical action (e.g., "approve"). A JavaScript helper library can be implemented on the client-side to perform this mapping. This library would subscribe to the `HostInterfaceBus` and monitor a stream of events over a short time window, perhaps 500 milliseconds. Within this window, it would analyze the patterns of activity—for example, the amplitude of sEMG signals from a specific muscle group, the direction of head tilt from an IMU, or the frequency of blinks detected by an eye-tracker. Using either a simple rule-based system or a lightweight machine learning model, the library would classify this pattern into one of a predefined set of `Intent` labels. Examples of such labels include `SPEAK`, `STOP`, `YES`, `NO`, `OVERLOAD`, `REQUEST_EXPLANATION_SLOW`, `SNAPSHOT`, and `RECOVERY_MODE`. These intents represent the core vocabulary of commands and states that the user wishes to express. Once an intent is confidently identified, the helper library would then package it into an RPC (Remote Procedure Call) request and send it to the host-node for execution. This design is reminiscent of systems like MyoGestic, which also focuses on mapping sEMG signals to discrete motor commands, but extends it to a broader range of modalities and semantic levels [26](#).

To further enhance reliability and reduce false positives, the system should incorporate mechanisms for error handling and uncertainty management directly at the bus level. For example, if the confidence score for a given `BciEvent` falls below a certain threshold, the event could be discarded or marked for secondary review. Similarly, the intent

classification logic in the JS helper should not operate on a single data point but on a consensus of evidence over a brief period. If multiple, independent sensors corroborate a single intent—for instance, a head nod (motion) and a specific EMG pattern (muscle contraction)—the system can assign a higher confidence to that intent before transmitting it. This multi-modal validation strengthens the integrity of the interface. Furthermore, the system must account for the possibility of no-intent. Periods of silence or neutral posture should not trigger unintended actions. The bus can be designed to recognize these null states and suppress event propagation until a clear, high-confidence intent is detected. This principle of fail-safe defaults is essential for preventing unwanted or erroneous interactions, especially in high-stakes situations .

The table below outlines a proposed schema for the `HostInterfaceBus` in Rust, including definitions for the `BciEvent` and `Intent` types, along with examples of standardized channel identifiers.

Component	Field / Variant	Type	Description
BciEvent Struct	timestamp_ns	u64	High-resolution timestamp in nanoseconds since Unix epoch. Enables cross-sensor synchronization.
BciEvent Struct	source_id	String	Unique identifier for the sensor device (e.g., "EMG_Bracer_Left", "EyeTracker_Muse_S2").
BciEvent Struct	channel_id	String	Standardized identifier for the measured feature (e.g., "EEG_THETA_POWER", "PUPIL_DIA_MM", "HRV_RMSSD"). See 19 for a full list.
BciEvent Struct	value	f64	Normalized, processed signal value.
BciEvent Struct	confidence	f32	Prediction or measurement confidence score (0.0 to 1.0). Lower values may be filtered out.
Intent Enum	SPEAK	()	Request to initiate speech synthesis or text-to-speech.
Intent Enum	STOP	()	Command to halt the currently active operation or output.
Intent Enum	YES	()	Affirmative consent or positive response.
Intent Enum	NO	()	Negative response or cancellation of an action.
Intent Enum	OVERLOAD	()	Signal that the user is experiencing cognitive or physiological overload. Triggers safety protocols.
Intent Enum	REQUEST_EXPLANATION_SLOW	()	Request for slower-paced explanations or simpler language.
Intent Enum	REQUEST_EXPLANATION_COMPLEX	()	Request for more detailed, complex, or technical information.
Intent Enum	SNAPSHOT	()	Capture a snapshot of the current state or context for later recall.
Intent Enum	RECOVERY_MODE	()	Enter a low-cognitive-load mode, simplifying UI and interactions.
Intent Enum	NO_INTENT	()	No discernible user intent detected during the observation window.

This standardized bus architecture provides the necessary foundation for building a robust and adaptable assistive system. It ensures that the host-node receives clean, structured, and semantically meaningful data, allowing its internal logic to focus on the higher-order tasks of authorization, adjustment, and adaptation, rather than getting bogged down in the complexities of low-level signal acquisition and parsing. This separation of concerns is key to creating a system that is both powerful and maintainable.

Mapping Intent to Action: The Six-Token Biophysical Authorship Model

Once the HostInterfaceBus has translated raw biological signals into high-level semantic Intents, the next critical step is to map those intents to quantifiable, safe, and authorized actions. This is accomplished through the six-token biophysical authorship model, a conceptual framework that represents the user's resources and capabilities as a bounded economy of six distinct tokens: BRAIN, WAVE, BLOOD, OXYGEN, NANO, and SMART . This model operationalizes the principle that every action, even a thought expressed through an assistive device, consumes a finite amount of biophysical energy and capacity. By framing these constraints as tokens, the system can enforce a strict accounting mechanism, ensuring that no operation is executed unless there is sufficient resource availability, thus protecting the user from cognitive and physiological overload . This approach transforms abstract neurorights into concrete, computable invariants that govern all interactions with the assistive system.

The first two tokens, BRAIN and WAVE, represent the user's cognitive and neural resources. They are analogous to a mental battery, measuring the expenditure of neural energy during cognitive tasks. These tokens are not direct measurements but are inferred from a fusion of multiple physiological signals. Research shows that increased mental workload is indexed by decreases in alpha power (8–12 Hz) and increases in theta power (4–7 Hz) in EEG signals, as well as by changes in spectral ratios like the frontal theta to parietal alpha ratio [21](#) [22](#) . Over-arousal states indicated by workload cause an incremental increase in heart rate and a decrease in Heart Rate Variability (HRV) metrics, providing a complementary measure of sympathetic nervous system activation [30](#) [42](#) . The WAVE token specifically captures dynamic changes in this neural load, while BRAIN could represent a longer-term baseline or total available capacity. Critically, any assistive action must be designed to consume only a small, bounded, and predictable amount of these tokens. For example, the SPEAK intent might correspond to a SystemAdjustment of WaveLoad { delta: -0.5 }, indicating a minor expenditure of wave energy, while an intent like REQUEST_EXPLANATION_COMPLEX would incur a much larger negative delta .

The next two tokens, BLOOD and OXYGEN, serve as vital sign indicators, representing the stability of the user's core physiological state. These tokens are designed to be protected from change by most external operations, ensuring that assistive functions never interfere with fundamental life-sustaining processes . Their values are derived from hemodynamic measurements obtained via technologies like functional Near-Infrared Spectroscopy (fNIRS), which measures concentrations of oxygenated hemoglobin (HbO),

deoxygenated hemoglobin (HbR), and total blood (HbT) in the brain [21](#) [38](#). Changes in these metrics correlate with brain activity and arousal states [42](#). Similarly, photoplethysmography (PPG) and electrocardiogram (ECG) sensors can provide continuous data on blood flow and oxygen saturation [20](#). In the biophysical model, **BL00D** and **OXYGEN** are treated as extremely sensitive parameters. An attempt by an external application or BCI to modify these tokens would be immediately rejected by the host-node's runtime, acting as a hard safety guardrail. This enforces the principle of non-maleficence, ensuring that the system enhances capability without compromising health .

The final two tokens, **NANO** and **SMART**, represent higher-order aspects of system governance and evolution, managed entirely by the user's inner ledger . These are not direct physiological measures but rather permissions and budgets controlled by the user. The **NANO** token can be conceptualized as a budget for low-level, autonomous optimizations and micro-decisions made by the system on the user's behalf. For example, adjusting screen brightness based on ambient light or choosing a default setting from a previously approved list could spend a small amount of **NANO**. The **SMART** token represents the user's autonomy budget for more significant system modifications, upgrades, or evolutionary changes. Upgrading a decoder model, enabling a new device, or altering core system policies would require spending **SMART** tokens . The key constraint is that the evolution of the system, including the modification of its own governing logic, requires explicit self-consent from the user . This prevents the system from becoming a black box that evolves in ways the user does not understand or approve of, thereby protecting the user's identity and sense of self [23](#). All **SystemAdjustment** events must be validated against these tokens, ensuring that the user retains ultimate authority over their own cognitive and system state.

A **SystemAdjustment** is the formal representation of an action's impact on the token economy. Each **RuntimeEventKind** triggered by an **Intent** must specify a precise delta for each relevant token. This must be done carefully to ensure safety and predictability. For example, an **AssistiveEvent :: Summarize** might only affect **WAVE** and **SMART**, but a hypothetical medical stimulation device would need to explicitly declare its effects on **BL00D** and **OXYGEN** to be considered compatible. The host-node's runtime enforces these adjustments, refusing any transaction that would violate the defined invariants, such as reducing **BL00D** or **OXYGEN** below a critical threshold or exceeding the user's available **WAVE** budget. This token-based accounting turns the abstract concept of "biophysical compatibility" into a hard-coded requirement for all integrations, forcing vendors and developers to adapt their systems to the user's sovereignty model rather than the other way around . The table below illustrates how

different **Intents** can be mapped to specific **SystemAdjustments** against the six tokens.

Intent	SystemAdjustment Example	BRAIN Change	WAVE Change	BLOOD/OXYGEN Change	NANO/SMART Change	Rationale & Safety Constraints
SPEAK	WaveLoad { delta: -0.3 }	None	-0.3	None	None	Represents minimal cognitive effort for basic communication. Zero impact on vital signs ensures safety .
NO	NoAdjustment	None	0.0	None	None	A simple negation that consumes negligible cognitive resources .
OVERLOAD	LifeforceHardStop	None	None	None	None	A signal to freeze all non-critical operations. Does not alter tokens but triggers system-wide safety protocols .
REQUEST_EXPLANATION_COMPLEX	WaveLoad { delta: -1.5 }, SmartAutonomy { delta: -0.2 }	None	-1.5	None	-0.2	Complex queries require a significant cognitive load (WAVE) and may necessitate a temporary reduction in system autonomy (SMART) to prevent system drift .
SNAPSHOT	NanoSpending { amount: 0.1 }	None	None	None	-0.1	Capturing a state is a low-level administrative task, consuming a small amount of the nano-autonomy budget .
EvolutionUpgrade(DecoderV2)	SmartAutonomy { delta: -1.0 }	None	None	None	-1.0	A major system upgrade requires a substantial expenditure of the smart autonomy budget, reflecting the need for conscious approval .

By implementing this six-token model, the system creates a transparent and auditable ledger of its own cognitive and physiological costs. Every interaction leaves a traceable record of its resource consumption, empowering the user to understand the true cost of their digital engagement. This model is a direct implementation of neurorights principles, particularly mental privacy and cognitive liberty, by ensuring that the user's cognitive resources are respected and not exploited by external entities [②](#). It provides a concrete, technical means of upholding the user's status as an "augmented citizen" whose authorship originates from within their own sovereign biological system .

Implementing the Core Runtime in Rust

The practical realization of the biophysical authorship model hinges on a robust and secure core runtime, primarily written in Rust for its memory safety and performance characteristics. This runtime acts as the central nervous system of the host-node, responsible for receiving events from the `HostInterfaceBus`, validating `SystemAdjustments` against the six-token economy, enforcing safety invariants, and maintaining an immutable log of all actions. The implementation must prioritize correctness, security, and adherence to the principles of host-locality, ensuring that no external entity can bypass its checks or compromise the integrity of the user's biophysical state. This involves defining key data structures, implementing a state transition function, and integrating with the underlying ledger and sensor drivers.

The first step in the Rust implementation is to define the core data structures that represent the system's state and transactions. This includes the `BciEvent` struct for incoming signals, the `RuntimeEventKind` enum for valid actions, and the `BioTokenState` struct to track the balance of the six tokens. The `BioTokenState` would contain six `f64` fields, one for each token (`brain`, `wave`, `blood`, `oxygen`, `nano`, `smart`), initialized to their maximum capacity upon system boot. The `RuntimeEventKind` enum defines all possible operations that can be performed, with each variant carrying the necessary parameters for its corresponding `SystemAdjustment`. For example, variants like `WaveLoad { delta: f64 }`, `SmartAutonomy { delta: f64 }`, and `NanoSpending { amount: f64 }` would be defined. This strongly-typed approach ensures that only legitimate operations can be constructed, preventing invalid state transitions at compile time.

```
// Example Rust code for core data structures

/// A standardized event from a biological sensor.
#[derive(Debug, Clone, Serialize, Deserialize)]
pub struct BciEvent {
    pub timestamp_ns: u64,
    pub source_id: String,
    pub channel_id: String,
    pub value: f64,
    pub confidence: f32,
}

/// The set of all valid actions the system can perform.
#[derive(Debug, Clone, Serialize, Deserialize)]
```

```

pub enum RuntimeEventKind {
    // Cognitive and Neural Load Adjustments
    WaveLoad { delta: f64 },
    BrainRecharge { delta: f64 },
    // Autonomy and Governance Adjustments
    SmartAutonomy { delta: f64 },
    NanoSpending { amount: f64 },
    // System State Management
    LifeforceHardStop,
    EvolutionUpgrade { details: String },
    // Assistive Actions (which trigger adjustments)
    AssistiveEvent { action_type: AssistiveAction },
}

/// The six biophysical tokens representing the user's state.
#[derive(Debug, Clone, Serialize, Deserialize)]
pub struct BioTokenState {
    pub brain: f64, // Total cognitive capacity
    pub wave: f64, // Current neural energy expenditure
    pub blood: f64, // Vital sign stability
    pub oxygen: f64, // Vital sign stability
    pub nano: f64, // Low-level autonomy budget
    pub smart: f64, // High-level system autonomy budget
}

```

With the data structures defined, the core of the runtime is the `apply_event` function. This function takes a `RuntimeEventKind` and the current `BioTokenState` as input and returns a new `BioTokenState` and a boolean indicating success or failure. This function is the sole gatekeeper for state changes. It must first check if the requested event is permissible given the current `LifeforceBand` and `SafetyCurveWave` invariants . For example, if the `wave` token has fallen into the `HardStop` band, the function should immediately reject any `WaveLoad` events. Second, it must validate the `delta` values for `SystemAdjustments`. For instance, it would verify that a `WaveLoad` operation does not exceed the `WAVE` ceiling and that it does not attempt to modify `BLOOD` or `OXYGEN`. Third, it must check the `SMART` and `NANO` budgets for any governance-related actions. If the `smart` budget is insufficient for an `EvolutionUpgrade`, the function fails. Only after passing all these checks is the state transition applied. This atomic, functional approach—where a new state is computed and returned instead of mutating the old one—prevents race conditions and makes the system's behavior deterministic and easier to reason about.

```

// Example Rust pseudo-code for the apply_event function

impl BioTokenState {
    pub fn apply_event(
        &self,
        event: &RuntimeEventKind,
        lifeforce_bands: &LifeForceBands,
        safety_curves: &SafetyCurves,
    ) -> Result<BioTokenState, &'static str> {
        // 1. Check LifeForce Bands and Safety Curves
        if lifeforce_bands.is_in_hard_stop(self.wave) {
            return Err("System in HardStop state. No WaveLoad events permitted");
        }

        if let RuntimeEventKind::WaveLoad { delta } = event {
            if safety_curves.is_overload(self.wave + delta) {
                return Err("Operation would exceed cognitive overload threshold");
            }
        }
    }

    // 2. Validate SystemAdjustments based on token rules
    let mut new_state = self.clone();
    match event {
        RuntimeEventKind::WaveLoad { delta } => {
            if delta.is_sign_negative() || *delta > self.wave {
                return Err("Invalid WaveLoad delta.");
            }
            new_state.wave -= delta;
        },
        RuntimeEventKind::BrainRecharge { delta } => {
            if delta.is_sign_negative() || (self.brain + delta) > MAX_BRAIN {
                return Err("Invalid BrainRecharge delta.");
            }
            new_state.brain += delta;
        },
        // ... other variants with similar validation logic
        RuntimeEventKind::LifeForceHardStop => {
            // Special event that doesn't change tokens but puts system in hard stop
            // Logic here depends on overall system design
        },
    }
}

```

```

        RuntimeEventKind::EvolutionUpgrade { .. } => {
            if self.smart < UPGRADE_COST {
                return Err("Insufficient SMART budget for evolution.")
            }
            new_state.smart -= UPGRADE_COST;
        },
        // ... other variants
    }

    // 3. Final sanity check to ensure no forbidden modifications
    if new_state.blood > self.blood || new_state.oxygen > self.oxygen
        return Err("Attempt to increase vital signs BLOOD/OXYGEN is fo")
    }

    Ok(new_state)
}
}

```

Finally, the runtime must integrate with the host-node's JSON-RPC server and the inner ledger. When an RPC call is received, the server parses the request, which should contain a `RuntimeEventKind`. The request must also include a `RpcSecurityHeader` containing the user's ALN/Bostrom identity and a cryptographic signature from their host-local keys, proving the request originated from the user's sovereign system . The runtime validates this signature and the event itself using the `apply_event` function. If successful, the new `BioTokenState` is committed to the ledger as a transaction, and the result is sent back to the client. This process ensures that all mutations to the system's state are signed, authorized, and recorded immutably. Adherence to standards like ISO 10993 for biocompatibility of any physical sensors used is also a prerequisite for collecting reliable data for this system [53](#) [58](#) . By building this logic in Rust, the system leverages the language's strengths to create a secure, performant, and verifiable core that faithfully implements the user's vision of biophysical authorship.

Developing the Client-Side Interface in JavaScript

While the Rust backend forms the secure and authoritative core of the system, the JavaScript client-side interface is the crucial bridge that allows external applications—such as AI chat interfaces, web browsers, and custom tools—to interact with the user's

biophysical state safely and effectively. This client-side component must be designed with several key responsibilities in mind: subscribing to the **HostInterfaceBus** to receive raw biosignals, classifying these signals into high-level **Intents**, packaging these **Intents** into budgeted RPC calls, and rendering a simplified, user-friendly feedback loop. The design must prioritize security by ensuring that no raw neural data ever leaves the local environment and that all communications are mediated through the host-node's authorization logic . This client-side logic effectively translates the user's subtle biological cues into deliberate, system-enforced commands.

The first responsibility of the JavaScript helper is to establish a connection to the **HostInterfaceBus**. This could be achieved through a WebSocket connection to the host-node's event streaming endpoint or via a browser-native API if the host-node exposes one. Upon connection, the helper subscribes to the desired **channel_ids** based on the user's configured preferences—for example, **EMG_RMS_CHANNEL**, **HEAD_TILT_DEGREES**, and **EEG_ALPHA_POWER**. It then begins to accumulate **BciEvent** objects into a sliding window buffer. This buffer serves as the raw material for the intent classification engine. To manage data volume and latency, the helper should only transmit events with a **confidence** score above a certain threshold, discarding noisy or uncertain readings locally. This local filtering is a key aspect of preserving mental privacy, as it prevents low-quality data from being unnecessarily processed or transmitted [12](#) .

The core functionality of the JavaScript helper is its intent classifier. This module runs locally in the user's browser and analyzes the aggregated events in the buffer to produce a semantic **Intent**. For simplicity and to avoid introducing complex models that could introduce bias or latency, this classifier can initially be rule-based. For example:

- An `Intent::YES` could be triggered if there is sustained, high-amplitude activity in a predefined `EMG_APPROVE_CHANNEL` for more than 300ms.
- An `Intent::NO` could be triggered by a rapid head shake detected via the IMU, provided it meets a minimum velocity threshold.
- An `Intent::OVERLOAD` might be triggered if a combination of rising EDA (indicating stress), decreasing HRV (indicating strain), and prolonged fixation on a single point is detected over a 10-second window.

As the system matures, this rule-based engine can be replaced or augmented with a lightweight machine learning model trained on the user's own self-data. This aligns with the principle of human-machine co-learning, where the system adapts to the individual user's unique biomarkers [11](#) . The MyoGestic framework, for instance, successfully uses a CatBoost classifier tailored to each participant, demonstrating the power of

personalization [26](#). The JavaScript helper would be responsible for feeding preprocessed data from the user's logged sessions into such a model for training, and then using the trained model for real-time inference. The output of this classifier is always a single **Intent** variant from the predefined set, ensuring consistency.

Once an **Intent** is classified with sufficient confidence, the helper must prepare and send an RPC request to the host-node. Crucially, this request must be budgeted. Before sending the **Intent** (e.g., `AssistiveEvent { action_type: Summarize }`), the helper must first query the host-node to determine the current **WAVE** and **SMART** token balances. Based on this information, it can decide whether the requested action is permissible. For example, if the user requests a complex summary but their **WAVE** budget is critically low, the helper could automatically downgrade the request to a simpler one (e.g., `ExplainInSimplestTerms`) or display a warning to the user. The RPC payload would contain the **Intent** and a **timestamp**. This request would then be signed locally with the user's browser-held private keys before being sent. The host-node verifies the signature, applies the **SystemAdjustment** using its Rust runtime, and sends back the result. This round-trip mediates all potentially costly operations through the biophysical authorship model.

The final piece of the JavaScript client is the feedback mechanism. Since the user cannot directly see their internal **BioTokenState**, the system must provide a proxy representation of their cognitive state. This could be a simple, non-distracting visual indicator in the corner of the screen, showing the relative levels of **WAVE** and **LIFEFORCE**. When the user performs an action, the helper could also provide haptic feedback (e.g., a gentle vibration) to confirm that the command was registered by the host-node. This closed-loop feedback is essential for the user to develop a feel for their resource consumption and to learn how to interact with the system in a way that is sustainable. It empowers the user to participate in the self-adaptation loop, as they can observe the consequences of their actions on their perceived state of fatigue or cognitive load. This aligns with recommendations for creating intuitive controls and clear "off switches" in assistive technologies, ensuring the user always feels in control [31](#). By handling these responsibilities—event subscription, local classification, budget-aware RPC mediation, and user feedback—the JavaScript client becomes an intelligent and respectful intermediary, enabling rich interactions while strictly respecting the user's biophysical sovereignty.

Building the Adaptive Assistive Language Loop

The most sophisticated and impactful component of the biophysical authorship system is the adaptive assistive language loop. This service is designed to leverage the power of large language models (LLMs) to aid communication and cognition without inducing the very cognitive overload it aims to alleviate. The core challenge is to create a closed-loop system where the LLM's assistance is dynamically priced in terms of the user's **WAVE** and **SMART** tokens, and where the system continuously learns from the user's responses to optimize this trade-off. This requires a careful orchestration of budgeted AI calls, difficulty tagging, and a data-driven feedback mechanism that refines the system's understanding of the user's optimal cognitive load interval over time .

The architecture of the assistive language loop begins with a small Rust/JS service that acts as a specialized client for the host-node's RPC interface. When a user provides a rough prompt—a keyword, a sentence fragment, or a conceptual tag—the service first consults the host-node to check the current **WAVE** and **SMART** token balances . Based on these balances, it determines a "budget" for the LLM call. This budget is not just a limit on cost but also on cognitive complexity. For instance, if **WAVE** is high, the service might request a long, complex response. If **WAVE** is low, it will constrain the LLM to generate a shorter, simpler response with a limited number of tokens. The prompt sent to the LLM is also templated to be safe and efficient, avoiding ambiguous instructions and focusing on concrete, single-step tasks . This ensures that even under budget constraints, the LLM's output remains useful and interpretable.

Upon receiving the LLM's output, the service performs two key analyses. First, it tags the output with a "difficulty level," which could be **simple**, **normal**, or **complex**. This tagging can be done by the LLM itself ("Please tag the following text as simple, normal, or complex:") or by a separate, lightweight classifier model trained on linguistic features known to correlate with readability, such as lexical density, sentence length, and syntactic complexity ¹⁹ . Second, the service mints a **SystemAdjustment** that reflects the cognitive cost of generating and processing that particular response. This adjustment, likely a **WaveLoad** event, is sent to the host-node and applied to the user's **BioTokenState**. This step is critical because it holds the system accountable for the cognitive debt it incurs, making the cost of assistance a tangible and measurable part of the user's experience .

The true intelligence of the loop comes from its self-refinement mechanism. The system must log every interaction: the initial prompt, the difficulty tag, the final output, the **WAVE** budget, the actual **WAVE** spent, and, most importantly, the user's subjective

feedback on the outcome (e.g., "too simple," "just right," "overwhelmed") . This logged data forms the basis for a personal adaptation pipeline. The user can periodically analyze this data to find correlations between their assisted interactions and their physiological state. For example, they might discover that making more than three complex requests per hour consistently leads to a drop in their **LifeForceBand** into the **SoftWarn** zone. This empirical finding can then be used to adjust the system's heuristics. The safety curves could be tuned to become more aggressive at lower **WAVE** levels, or the default "simple mode" threshold could be lowered .

Advanced machine learning techniques can further enhance this personalization. Research shows that linear mixed-effects models can effectively identify individual predictors of cognitive load, such as specific EDA and HRV markers ³⁰ . Applying such a model to the logged data could allow the system to move beyond fixed thresholds and toward a personalized "optimal load interval." Instead of a binary **HardStop** vs. **Safe** state, the system could operate on a continuum, using fuzzy or neuro-fuzzy logic to smoothly adjust the difficulty of outputs based on a probabilistic assessment of the user's current state ²² . This acknowledges that cognitive states are dynamic and non-binary, and a one-size-fits-all approach to overload detection is suboptimal ²² . The goal is to create a symbiotic relationship where the AI assists the user, and the user's biological feedback helps the AI better tailor its assistance. This approach directly addresses the risks of cognitive offloading and dependency identified in studies on AI reliance, by making the user an active participant in managing their cognitive engagement ²³ . The end result is an assistive loop that is not just reactive but truly adaptive, evolving alongside the user to provide support precisely when and how it is needed, without ever compromising their cognitive well-being.

Enforcing Security, Governance, and Neurorights Compliance

The entire biophysical authorship architecture is built upon a foundation of stringent security and governance principles, designed to translate abstract neurorights into concrete, enforceable technical and contractual obligations. The central tenet is the absolute primacy of host-local authority: all signing and authorization decisions must originate from the user's personal hardware and cryptographic keys, never from a remote cloud service . This design choice directly counters the primary attack vectors of account takeover and coercive monitoring. Every interaction, from a simple YES intent to a major

system evolution, must pass through the user's sovereign inner ledger, which maintains an immutable, tamper-evident audit trail of all events . This ledger serves as the ultimate arbiter of legitimacy and accountability.

To enforce this model on third-party integrations, the user can embed governance clauses directly into their ALN/Bostrom identity shard. These are not mere suggestions but hard-coded requirements that any device or service must satisfy to be granted access. The policy templates should mandate strict adherence to the `HostInterfaceBus` schema, prohibiting any deviation that would require bespoke client-side handling. More importantly, they must stipulate that all effects on the biophysical token economy must be expressed as safe `SystemAdjustment` events, fully subject to the protection of the host-node's `LifeforceBand` and `SafetyCurveWave` invariants . A sample clause for such a governance shard might read: "Any integration must map its functionalities exclusively to the `RuntimeEventKind` enumeration defined in the host-node's protocol. It may not propose or execute any operation that modifies the `BL00D` or `OXYGEN` tokens, nor may it bypass the host-node's state validation logic." This transforms "biophysical compatibility" from a feature into a non-negotiable requirement, effectively forcing a vendor-neutral, user-centric standard .

These technical safeguards are designed to operationalize the five core neurorights discussed in legal and ethical frameworks ² . **Mental Privacy** is protected by mandating on-device processing and minimal data export; raw neural data never leaves the host, and only purpose-bound summaries are ever exposed ^{12 31} . **Cognitive Liberty** is upheld by ensuring that token budgets are determined by the user's biological state, not by subscriptions, financial incentives, or vendor-imposed limits ³⁰ . **Mental Integrity and Identity** are preserved by placing hard guards around the `BL00D` and `OXYGEN` tokens, preventing any external agent from manipulating the user's core physiological state, and by requiring explicit self-consent for any system evolution that could alter the user's cognitive profile . **Fair Access** is guaranteed by the non-financial nature of the inner ledger, which treats all eligible roles (e.g., `AugmentedCitizen`, `AuthorizedResearcher`) equally, regardless of economic status . Finally, the right to **Biophysical Authorship** is the very foundation of the system, ensuring that the user's biology and host-local keys are the sole validators of actions and signatures, and that no platform can act as a super-user or override the user's will .

The enforcement of these policies relies on a multi-layered security architecture. The host-node's RPC endpoints must be fortified with strict identity and permission checks. Every incoming request must carry a `RpcSecurityHeader` containing the sender's identity and a valid cryptographic signature from a key allowed by the governance shard . The system must distinguish between different roles, such as a generic `Vendor` role and

a more privileged `AssistiveTool` role, and grant permissions accordingly. Sandbox environments, for example, should be denied the ability to submit mutating `RuntimeEventKind` events. Furthermore, all data exchanged between the client and the host-node, and all entries in the inner ledger, must be encrypted to protect against eavesdropping and data breaches ¹. By combining these technical controls—host-local signing, immutable ledgers, budgeted RPCs, granular role-based access control, and strong encryption—with clear, neurorights-aligned governance policies, the user can construct a comprehensive defense. This holistic approach ensures that as they augment their capabilities with technology, they do not cede their fundamental rights and freedoms, but instead build a robust, personalized infrastructure for exercising them.

Reference

1. Biomedical applications of wearable biosensors <https://www.sciencedirect.com/science/article/pii/S2949822823000849>
2. (PDF) Brain-MCP: A Fully Homomorphic Neuro-Symbolic ... https://www.researchgate.net/publication/393129669_Brain-MCP_A_Fully_Homomorphic_Neuro-Symbolic_Protocol_for_Governing_System_3_Cognition_and_Restoring_Semantic_Scarcity
3. TinyML for Ultra-Low Power AI and Large Scale IoT ... <https://www.mdpi.com/1999-5903/14/12/363>
4. national conference on engineering innovations in ... <https://ijsrst.com/paper/v12i15.pdf>
5. Applications of Machine Learning in Assessing Cognitive ... https://www.researchgate.net/publication/397226464_Applications_of_Machine_Learning_in_Assessing_Cognitive_Load_of_Uncrowded_Aerial_System_Operators_and_in_Enhancing_Training_A_Systematic_Review
6. Machine Learning Apr 2025 <https://www.arxiv.org/list/cs.LG/2025-04?skip=1525&show=2000>
7. TechRxiv - - Advanced Genetics Preprints - https://ggnpreprints.authorea.com/inst/26407?page=156&sortby=Most+Cited&tag_filter=Signal+Processing+And+Analysis
8. A Survey on Haptics: Communication, Sensing and ... <https://ieeexplore.ieee.org/iel8/9739/5451756/10637258.pdf>

9. Technological enhancements in personalized dietary ... <https://www.sciencedirect.com/science/article/pii/S2667099225000374>
10. Network Challenges for Cyber Physical Systems with Tiny ... <https://www.mdpi.com/1424-8220/15/4/7172>
11. Human-Machine Co-Learning: interactive curriculum ... https://theses.hal.science/tel-04828514/file/142648_SUNGEELLE_2024_archivage.pdf
12. Rust Implementation of the Privacy-Preserving Funshade ... <https://dl.acm.org/doi/full/10.1145/3733802.3764061>
13. Deep learning applied to EEG data with different montages ... https://www.researchgate.net/publication/377516929_Deep_learning_applied_to_EEG_data_with_different_montages_using_spatial_attention
14. Arxiv今日论文 | 2025-11-03 http://lonepatient.top/2025/11/03/arxiv_papers_2025-11-03
15. AI for Good Global Summit 2024 - ITU <https://aiforgood.itu.int/summit24/>
16. Current 20251002 | PDF | Facebook | Obscenity <https://www.scribd.com/document/929887999/CURRENT-20251002>
17. Belli & Gaspar (Eds.) - 2023 - The Quest Sovereignty ... <https://www.scribd.com/document/677637169/Belli-Gaspar-eds-2023-The-Quest-Sovereignty-Transparency-Accountability>
18. 3D Printing Assisted Wearable and Implantable Biosensors <https://pmc.ncbi.nlm.nih.gov/articles/PMC12468503/>
19. A Survey on Measuring Cognitive Workload in Human ... <https://dl.acm.org/doi/10.1145/3582272>
20. Human-Centric Cognitive State Recognition Using ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12252513/>
21. Mental workload assessment by monitoring brain, heart ... <https://www.frontiersin.org/journals/neuroergonomics/articles/10.3389/fnrgo.2024.1345507/full>
22. Assessing Cognitive Load Using EEG and Eye-Tracking in ... <https://www.mdpi.com/2414-4088/9/9/99>
23. Measuring Visual Fatigue and Cognitive Load via Eye ... https://www.researchgate.net/publication/354947961_Measuring_Visual_Fatigue_and_Cognitive_Load_via_Eye_Tracking_while_Learning_with_Virtual_Reality_Head-Mounted_Displays_A_Review
24. Computer Science <https://www.arxiv.org/list/cs/new?skip=25&show=1000>
25. Latest articles https://www.mdpi.com/latest_articles

26. MyoGestic: EMG interfacing framework for decoding ... <https://www.science.org/doi/10.1126/sciadv.ads9150>
27. Mapping EEG Metrics to Human Affective and Cognitive Models <https://pmc.ncbi.nlm.nih.gov/articles/PMC12649996/>
28. (PDF) A Preliminary Assessment of Neuro-Salutogenic ... https://www.academia.edu/81944358/A_Preliminary_Assessment_of_Neuro_Salutogenic_Landscape_Stimuli_in_Neighbourhood_Parks_Theory_Based_Model_for_Stress_Mitigation
29. Mepco r2023 Cse Bcs | PDF | Engineering <https://www.scribd.com/document/751674354/Mepco-r2023-Cse-Bcs>
30. Detection of Cognitive Load Modulation by EDA and HRV - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC12030397/>
31. (PDF) Exploring Eye Tracking to Detect Cognitive Load in ... https://www.researchgate.net/publication/386368327_Exploring_Eye_Tracking_to_Detect_Cognitive_Load_in_Complex_Virtual_Reality_Training
32. IoT-Based Approaches to Personnel Health Monitoring in ... <https://www.mdpi.com/2071-1050/18/1/365>
33. Ethical Aspects of Cyber-Physical Systems [https://www.europarl.europa.eu/RegData/etudes/STUD/2016/563501/EPRS_STU\(2016\)563501\(ANN1\)_EN.pdf](https://www.europarl.europa.eu/RegData/etudes/STUD/2016/563501/EPRS_STU(2016)563501(ANN1)_EN.pdf)
34. (PDF) Neural Efficiency and Sensorimotor Adaptations in ... https://www.researchgate.net/publication/400101755_Neural_Efficiency_and_Sensorimotor_Adaptations_in_Swimming_Athletes_A_Systematic_Review_of_Neuroimaging_and_Cognitive-Behavioral_Evidence_for_Performance_and_Wellbeing
35. Program in Chronological Order <https://ieeexplore.ieee.org/iel5/4352184/4352185/04352201.pdf>
36. Changelog — BIND 9 9.21.15 documentation <https://bind9.readthedocs.io/en/v9.21.15/changelog.html>
37. An adaptive hand exoskeleton rehabilitation training ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC12728062/>
38. Mental workload assessment by monitoring brain, heart, and ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC10963413/>
39. sitemap-questions-51.xml <https://stackoverflow.com/sitemap-questions-51.xml>
40. WISE 2024 PhD Symposium, Demos and Workshops <https://link.springer.com/content/pdf/10.1007/978-981-96-1483-7.pdf>
41. (PDF) Wearable EEG-Based Cognitive Load Classification ... https://www.researchgate.net/publication/369015746_Wearable_EEG-

Based_Cognitive_Load_Classification_by_Personalized_and_Generalized_Model_Using_Brain_Asymmetry

42. A Review of Psychophysiological Measures to Assess ... <https://www.frontiersin.org/journals/human-neuroscience/articles/10.3389/fnhum.2019.00057/full>
43. An Overview of Stress Analysis Based on Physiological ... <https://www.mdpi.com/1424-8220/25/23/7108>
44. International Conference for Innovation in Biomedical ... <https://link.springer.com/content/pdf/10.1007/978-981-10-0266-3.pdf>
45. SLEEP | Volume 36 | Abstract Supplement <https://academic.oup.com/DocumentLibrary/SLEEP/2013abstractSupplement.pdf>
46. Recent Trends of AI Technologies and Virtual Reality <https://link.springer.com/content/pdf/10.1007/978-981-96-1154-6.pdf>
47. Smart and Sustainable Intelligent Systems https://www.researchgate.net/profile/Deepika-Kukreja/publication/350376595_A_Means_of_Futuristic_Communication_A_Review/links/6530bc841d6e8a70703c827d/A-Means-of-Futuristic-Communication-A-Review.pdf
48. sitemap-questions-157.xml <https://stackoverflow.com/sitemap-questions-157.xml>
49. Civ Bci R19 | PDF | Internet Of Things | Smart Grid <https://www.scribd.com/document/906025539/CIV-BCI-R19>
50. Linknovate | Profile for Stack Overflow <https://www.linknovate.com/affiliation/stack-overflow-1884317/all/?query=unique%20row%20id>
51. (PDF) Switching on Smart Care: The Ascendancy of ... https://www.researchgate.net/publication/397370882_Switching_on_Smart_Care_The_Ascendancy_of_Wireless_Technologies_in_Continuous_Health_Surveillance
52. Program in chronological order <http://ieeexplore.ieee.org/iel7/7580725/7590615/07590623.pdf>
53. Rubber Carbon Electrodes - Durable and Versatile Solutions <https://www.alibaba.com/showroom/rubber-carbon-electrodes.html>
54. Silver Carbon Fiber Fabric: Real-World Performance ... <https://www.aliexpress.com/p/wiki/article.html?keywords=silver-carbon-fiber-fabric>
55. Conductive Silicone Rubber Electrodes for Tens & Ems <https://www.alibaba.com/showroom/conductive-silicone-rubber-electrode.html>
56. 张威SSAT类比 <https://www.scribd.com/document/806092800/%E5%BC%A0%E5%A8%81SSAT%E7%B1%BB%E6%AF%94>
57. Multi-Channel Soft Dry Electrodes for Electrocardiography ... <https://www.mdpi.com/1424-8220/24/2/420>

58. A multifunctional platform for the production and ... <https://www.sciencedirect.com/science/article/pii/S0924424725002973>
59. PulmoVista 500 <https://www.draeger.com/Content/Documents/Content/pulmovista-500-sw-12n-ifu-9055985-en.pdf>
60. ISO 10993-1:2025(en), Biological evaluation of medical ... <https://www.iso.org/obp/ui/en/#!iso:std:84512:en>
61. ISO 10993-1 and Biocompatibility https://www.emergobyul.com/sites/default/files/2023-08/PLC23CS1070961-ISO-Biocompatibility-Whitepaper_Final-Digital.pdf
62. INTERNATIONAL STANDARD ISO 10993-1 [https://mdcpp.com/doc/standard/ISO10993-1-2018\(E\).pdf](https://mdcpp.com/doc/standard/ISO10993-1-2018(E).pdf)
63. ISO 10993-1:2025 - Biological evaluation of medical devices <https://www.iso.org/standard/10993-1>
64. ISO 10993-1:2025 Revision Highlights Part 1 <https://medinstitute.com/blog/iso-10993-12025-revision-highlights-part-1-reasonably-foreseeable-misuse/>
65. Design Control Guidance For Medical Device Manufacturers <https://www.fda.gov/media/116573/download>
66. General Principles of Software Validation - Final Guidance ... <https://www.fda.gov/media/73141/download>
67. Design Controls <https://www.fda.gov/inspections-compliance-enforcement-and-criminal-investigations/inspection-guides/design-controls>
68. Design Controls <https://www.fda.gov/files/drugs/published/Design-Controls---Devices.pdf>
69. FDA Hosted PDF <https://manuals.plus/m/01f1509cb9f0615dbea3d68e2958a749ef7f262bb22e05eb48bcf095be9b187b>