Governing Self-Directed Augmentation: A Dual-Layered Blueprint for Neurorights, Ecosafety, and Voluntary Evolution

Policy Architecture: A Dual-Layered Constitutional System

The development of a fair and non-coercive cybernetic evolution policy necessitates a sophisticated architectural design capable of managing the inherent tensions between individual agency, systemic safety, and ecological accountability. The proposed solution is not a single, monolithic regulatory body but a dual-layered constitutional system. This architecture explicitly separates the concerns of human autonomy from those of environmental stability, creating two parallel yet interconnected governance layers. The first layer, the ecosafety spine, functions as a foundational, non-negotiable safety net. It operates independently to monitor and enforce ecological stability by consuming standardized metrics such as Knowledge (K), Eco-impact (E), and Risk-of-harm (R), while enforcing Lyapunov-style residual constraints ($V_{t+1} \leq V_t$) [37] [38] . Its primary mandate is to ensure that any technological or evolutionary path does not destabilize the broader socio-technical ecosystem. The second layer, the neurorights firewall, is a parallel, constitutionally-governed entity that sits atop the ecosafety spine. It consumes the safety outputs from the spine as inputs for its own decision-making but maintains a distinct logic based on principles of consent, fairness, and non-discrimination . This separation is critical; it allows the system to reject an evolution path as ecologically unsound without making a judgment on an individual's right to pursue it, thereby upholding the core principle of voluntary participation.

This dual-layer model provides a robust framework for balancing competing priorities. The ecosafety spine acts as a hard constraint on "how" an upgrade can be achieved, focusing on quantitative measures of safety and sustainability. For example, if a proposed neural interface design increases the Lyapunov residual beyond a safe threshold, the spine would flag it as a threat to systemic stability, regardless of its potential benefits to the user . Conversely, the neurorights firewall governs the question of "whether" an upgrade is permissible, focusing on qualitative aspects of human rights and dignity. It would protect an individual's decision to undergo an upgrade, even if the ecosafety spine identifies some level of acceptable risk, potentially triggering a higher-level review or requiring additional safeguards rather than an outright rejection . This structure aligns with international legal trends that employ risk-based approaches to regulate technology, such as the EU AI Act, which prohibits certain high-risk practices like real-time remote biometric identification and social scoring systems [14] [21]. By decoupling these concerns, the policy avoids the pitfalls of centralized control and creates a resilient system where neither ecological stability nor human freedom can be compromised in the name of the other. The concept of resilience itself, defined as an ecosystem's capacity to absorb disturbances and maintain its structure, serves as a powerful metaphor for this entire governance framework [43] [44].

The interaction between these two layers is carefully choreographed. The ecosafety spine generates a set of quantifiable safety metrics, including normalized risk coordinates $r_x \in [0,1]$ for various impact categories (toxicity, heat, etc.) and a global Lyapunov residual $V_t = \sum w_j r_j$ [37]. These values are then passed as inputs to the neurorights firewall. The firewall uses these metrics to inform its evaluation of an evolution request but does not override them. Instead, it applies its own set of constitutional rules. For instance, a proposed evolution path might satisfy the ecosafety spine's criteria (e.g., $V_{t+1} \leq V_t$) but fail the neurorights firewall's test if it involves a coercive consent mechanism or leads to neurodiscrimination against non-augmented individuals [12]. The final decision is therefore a composite one, ensuring that all upgrades are both ecologically sustainable and ethically sound. This design choice reflects a deep understanding of the challenges posed by emerging technologies, drawing parallels to the need for new legal frameworks to address neurotechnology, as discussed by organizations like the OECD and UNESCO [19] [53]. By treating the ecosafety spine as a shared, universal grammar that all future tools must compile against, the policy ensures that every assistive technology contributes to tightening Earth-healing corridors rather than operating in a separate, unaccountable layer . This makes the loop of technological creation and deployment a self-regulating, closed system focused on restoration and safety.

High-Level Policy Constraints: The Neurorights Constitution

The foundation of the cybernetic evolution policy rests upon a set of inviolable principles, forming a constitution that governs the relationship between the individual, the augmented citizen, and the system. These constraints are not mere suggestions but hardcoded limitations designed to ensure fairness, voluntariness, and the protection of fundamental human rights in the age of neurotechnology. The first and most critical principle is strict voluntariness and non-coercion. No individual can be forced to accept evolution upgrades, and refusal to evolve must never result in the denial of access to basic services, information, or fundamental rights . This principle is a direct application of the concept of "cognitive liberty," identified as a key ethical safeguard that protects an individual's ability to make free and competent decisions about the use of neurotechnologies for therapeutic or enhancement purposes [12] [19] . This aligns with international human rights law, which imposes both negative obligations (not to engage in prohibited acts like torture) and positive obligations (to take active measures to protect rights like privacy and freedom of thought) [14] . The policy must encode this principle through technical means, ensuring that any attempt to withhold essential services from a non-evolved individual constitutes a violation of the core governance rules.

A second cornerstone of the policy is the establishment of a hard ceiling on Risk-of-Harm (RoH). Any proposed evolution workflow, from initial design to final implementation, must project a RoH value of no more than 0.3 . If a projection exceeds this threshold, the system must automatically trigger a rejection or escalate the request for a higher-level review by a multi-stakeholder panel. This quantitative ceiling transforms abstract safety concerns into a clear, measurable boundary, enabling automated enforcement. The process of setting and monitoring this margin can draw inspiration from methodologies used in clinical trials, such as Bayesian non-inferiority approaches that incorporate expert opinions on acceptable risk levels [98] [100]. This approach moves beyond subjective ethical assessments and grounds safety in verifiable, data-driven models, similar to the adaptive stability learning methods applied to nonlinear dynamical systems [37] . The goal is to create a system where personal advancement is bounded by a scientifically-grounded measure of safety, preventing the pursuit of capabilities at an unacceptable cost to well-being.

The third major constraint is the explicit protection of non-augmented individuals from harm and discrimination. The policy must actively guard against "neurodiscrimination," which is defined as discrimination based on neural signatures that may indicate predispositions to dementia, mental health conditions, personality traits, cognitive performance, intentions, or emotional states [12] . This requires implementing affirmative protections, ensuring that a person's decision not to augment does not become a basis for social, economic, or legal disadvantage. This principle is supported by emerging legal frameworks, such as the Council of Europe's Convention on Artificial Intelligence, which obligates states to align their AI governance with human rights, democracy, and the rule of law [14] . Furthermore, the policy must uphold the right to mental integrity and identity, which includes the right to a sense of self and the right to free will [20] [65]. This implies that no system can score or profile an individual's inner state without explicit, revocable consent, a concept reinforced by regulations like GDPR, which allow for the withdrawal of consent for data processing [1]. Finally, the principle of revocability at will is paramount. Every augmentation must come with a guaranteed ability to reverse or remove it, serving as a crucial long-term safeguard against unforeseen consequences. This mirrors the importance of validation and verification processes in safety-critical systems, where the ability to return to a known safe state is a fundamental requirement [41]. Together, these four principles—voluntariness, RoH ceilings, anti-neurodiscrimination, and revocability—form a comprehensive constitutional framework that prioritizes human dignity and autonomy above all else.

Principle Constraint	Description	Supporting Legal/Conceptual Basis
Voluntary Participation	No coercion; refusal to evolve cannot lead to exclusion from basic services or rights.	Cognitive Liberty [19]; Negative Obligations in Human Rights Law [14]; EU AI Act (Anti-Social Scoring) [14]
Hard RoH Ceiling	All evolution workflows must project a Risk-of-Harm $\leq$ 0.3. Projections $>$ 0.3 trigger rejection or review.	Adaptive Control Theory [38]; Bayesian Non-Inferiority Testing [100]; Lyapunov Stability Analysis [37]
Protection of Non-Augmented	Prohibition of neurodiscrimination based on neural signatures or lack thereof.	Neurodiscrimination [12]; Right to Mental Integrity [19]; Council of Europe AI Convention [14]
Revocability at Will	Consent for an upgrade must be revocable, with a guaranteed technical pathway for rollback.	GDPR Data Protection [1]; Validation & Verification in Safety-Critical Systems [41]

Technical Interoperability Standards: The ALN-Shard Connector

To translate the high-level policy constraints into an operational reality, a common technical language and standardized data format are essential. The proposed solution is the adoption of a formal Algebraic Logic Notation (ALN) shard schema, specifically the `evolution.connector.v1` shard, which serves as the central artifact for governing every augmentation event. This shard acts as a digital envelope containing all necessary metadata, consent records, and safety metrics for a single evolution step, ensuring that every action is transparent, auditable, and compliant with the neurorights constitution . The design of this shard is minimalist yet comprehensive, containing several critical fields. First, it must include a `consent.envelope` field, which digitally binds the evolution request to the user's Decentralized Identifier (DID) and contains a record of explicit, logged, and revocable consent, satisfying the core principle of voluntariness [12]

[106]. Second, it must define the specifics of the upgrade through `capability.negotiation` fields, detailing the type (e.g., cognitive assist, sensorium) and intensity of the proposed change . Third, it must embed a `revocation.flag` field, which initiates the predefined rollback and undo pathways when an individual chooses to terminate an augmentation .

Beyond its data fields, the ALN shard schema must enforce a set of immutable invariants —automated checks that act as guards within the system itself. These invariants codify the policy directly into the code, making it computationally difficult to bypass core safety rules. Drawing from the patterns established in the ecosafety spine, several key invariants are necessary. The `no_corridor_no_build` invariant ensures that any variable being modified has a corresponding corridor definition, preventing actions on undefined parameters . The `safestep` invariant rejects any proposed evolution path that would cause a normalized risk coordinate r_x to fall outside its defined `[0,1]` bounds or increase the overall Lyapunov residual V_t, thus maintaining ecological stability . Finally, the `ker_delta` invariant mandates that each evolution step must result in a non-degrading outcome across the K/E/R triad, meaning knowledge and eco-impact cannot decrease while risk cannot increase beyond a specified threshold . These invariants provide a layered defense, checking everything from data completeness to systemic stability at the moment of execution. By defining these constraints in a machine-checkable format like ALN, the system can automatically validate every evolution request before it is processed, ensuring continuous adherence to the policy's core tenets .

The practical implementation of this technical standard begins with the development of Rust modules that parse and enforce the ALN schema. The provided context offers a template for such a system, demonstrating how a `ResponseShard` can be structured with fields for DID, topic, K/E/R scores, and a calculated residual, along with a method to verify improvements over a previous state . This can be directly adapted for the `evolution.connector.v1` shard. The following table outlines the proposed minimal schema for this connector.

Field Name	Type	Description
`user_did`	String	The user's Decentralized Identifier, linking the evolution event to a specific, verifiable agent [106].
`topic`	String	The subject area or domain of the evolution (e.g., "phoenix-mar-sat", "biodegradable-materials") .
`triad`	Struct	The K/E/R scores for the proposed evolution, representing its impact on knowledge, ecology, and safety .
`residual`	Struct	The Lyapunov-style residual V_t and its constituent risk coordinates, used to check systemic stability .
`evidence`	Vec	A log of supporting data, calculations, or external references justifying the evolution proposal .
`corridor_tags`	Vec	Tags referencing the specific corridor definitions (e.g., "mar", "sat", "phoenix") that apply to this evolution .
`revocation_flag`	Boolean	A flag indicating whether this shard is part of a revocation or rollback workflow .

By establishing this shard as the universal format for all augmentation events, the policy creates a common ground for all stakeholders—individuals, developers, and governance bodies—to interact with the system. It turns abstract principles into concrete, verifiable data points, enabling automated enforcement, transparent auditing, and reliable dispute resolution. This technical backbone is what gives the neurorights constitution its teeth, transforming it from a philosophical document into a functional, secure, and scalable governance framework.

Procedural Governance Frameworks: Consent, Revocation, and Arbitration

While technical standards provide the rigid structure for the cybernetic evolution policy, procedural governance frameworks supply the necessary flexibility to manage complex human interactions, conflicts, and changes. These procedures ensure that the policy is not only technically sound but also practically fair and accountable. A key procedural element is the implementation of multi-party consent protocols. While most evolution events require only the consent of the augmented citizen, upgrades deemed "powerful"—those with the potential to significantly influence others, control critical infrastructure, or alter governance flows—may necessitate stricter approval mechanisms . For such cases, consent could be required from multiple parties, including designated governance roles or community representatives, not just the individual seeking the upgrade. This prevents an individual from accumulating disproportionate power through their enhancements, acting as a crucial check against the emergence of coercive hierarchies. This concept is analogous to multi-signature wallets in blockchain systems, where transactions require approval from more than one private key, adding a layer of collective oversight.

Equally important is the establishment of guaranteed rollback and undo pathways. The principle of "revocable-at-will" must be technically enforceable, not merely a legal abstraction . Every evolution connector shard must contain a detailed, pre-approved plan for rolling back the specific changes it introduces. This includes reverting hardware configurations, deleting software implants, and restoring baseline cognitive or physiological states. The existence of this pathway is a prerequisite for any upgrade to be approved. This mirrors the rigorous validation and verification processes found in flight-critical operations, where the ability to safely abort a mission or return to a stable state is a primary design consideration [41] . The revocation flag within the ALN shard triggers this predefined workflow, ensuring that the promise of reversibility is honored. This provides a vital safety net, allowing individuals to reassess their choices and withdraw from an augmentation path without facing irreversible consequences.

In cases where disputes arise regarding the validity, impact, or revocation of an upgrade, a clear jurisdictional arbitration process is indispensable. This framework must define the rules for resolving conflicts between individuals, developers, and governance bodies. It should reference established legal precedents and international standards to ensure consistency and legitimacy. For example, rulings from the Chilean Supreme Court concerning brain data privacy offer valuable guidance on protecting mental integrity in the face of neurotechnology [55] [81] . Similarly, recommendations from the Organisation for Economic Co-operation and Development (OECD) on responsible innovation in neurotechnology provide a global benchmark for ethical conduct [19] [53] . The arbitration process should be transparent, with all decisions and their rationales recorded as auditable traces, potentially hex-stamped for immutability. Finally, the entire procedural framework must support a closed-loop feedback system for policy refinement. Every evolution step, along with its outcomes and any associated risks or impacts, should be logged and analyzed. This data, comprising the K/E/R scores and RoH projections, forms the basis for iterative improvement of the policies, caps, and connector designs themselves . This quantitative learning process ensures that the policy evolves over time, becoming more precise and effective at balancing progress with safety, always under the guiding light of the established principles.

Integrating Ecosafety: Quantifying Impact on Evolution Pathways

The integration of the cybernetic evolution policy with an overarching ecosafety spine is a defining feature of this governance model, moving beyond abstract ethical considerations to establish concrete, quantifiable constraints on human enhancement. This tight coupling ensures that individual advancement is never pursued at the expense of systemic stability. The evolution policy is designed to consume K/E/R (Knowledge/Eco-impact/Risk-of-harm) metrics and Lyapunov residual constraints generated by the ecosafety spine as mandatory inputs for every proposed evolution path . An upgrade that is deemed personally beneficial but carries an unacceptable ecological footprint (resulting in a low 'E' score) or destabilizing effect (causing the Lyapunov residual V_t to increase) would be rejected by the combined system. This creates a symbiotic relationship where technological progress is made contingent upon its contribution to environmental healing and sustainability. The assistive technologies developed to aid in this process are themselves governed by the same ecosafety spine, converting "assistive tech" from a vague booster into a rigorously scored optimizer of Earth-restoration corridors .

The practical application of this integration relies on a shared vocabulary of metrics and constraints. The ecosafety spine defines a set of normalized risk coordinates $r_x \in [0,1]$ for variables across different domains, such as toxicity, heat, hydrology, and social load [37]. These coordinates contribute to a global Lyapunov residual, $V_t = \sum w_j r_j$, which represents the system's overall deviation from a stable, safe state [37]. Any admissible move, including a proposed evolution, must satisfy the constraint $V_{t+1} \leq V_t$ [37]. The evolution policy leverages this framework by anchoring its evaluations to specific "Phoenix" nodes, which are concrete examples of Earth-healing modules like MAR cells, biodegradable materials, or circular hardware . For instance, a proposed cognitive enhancement for an engineer working on a Phoenix-class MAR cell pilot would be evaluated based on how that upgrade tightens the cell's specific corridor constraints, improves its normalization kernels, or lowers its residual risk score . This forces every new tool or capability to demonstrate a tangible, positive impact on a real-world ecological problem, effectively closing the loop between technological development and environmental accountability .

The concept of ecological resilience provides a powerful conceptual lens for this entire framework. Resilience is defined as an ecosystem's capacity to absorb disturbances and maintain or quickly recover its structure and function [43] [44]. The evolution policy can be understood as a mechanism to ensure that human cybernetic evolution does not push the socio-technical system past its own resilience threshold. The four-dimensional resilience index, based on the paradigm of resistance, recovery, reconstruction, and renewal, offers a sophisticated model for measuring the health and stability of the system being governed [42]. By framing the problem in terms of resilience, the policy acknowledges that stability is not a static state but a dynamic capacity to adapt. The Lyapunov residual, in this context, is a measure of the system's distance from a tipping point. The evolution policy, by enforcing the $V_{t+1} \leq V_t$ constraint, acts as a regulator, keeping the system within a safe, resilient operating space. This approach is consistent with modern ecological management strategies that focus on enhancing resilience to cope with uncertainty and shocks [45] [47]. Ultimately, the integration of ecosafety transforms the goal of cybernetic evolution from pure capability maximization to a more holistic objective: enhancing human potential in a way that simultaneously strengthens the resilience and health of the planet itself.

Implementation Roadmap: From Specification to Pilot Deployment

The development of the cybernetic evolution policy can be approached through a phased roadmap, beginning with the formalization of its core components and progressing toward the deployment of a concrete pilot project. The initial phase must focus on creating a detailed specification document that precisely defines the architecture of the dual-layered system. This document should delineate the interfaces between the neurorights firewall and the ecosafety spine, specifying exactly what data flows from Layer 1 to Layer 2 and what decisions are made autonomously by each layer. Key to this stage is the development of the `evolution.connector.v1` ALN shard schema. This involves designing the minimal set of fields required for governance, starting with the user's DID, topic, K/E/R triad, residual, evidence, and corridor tags, and implementing the core invariants (`safestep`, `ker_delta`) in a Rust library to validate their logical soundness . Concurrently, the high-level policy principles—voluntariness, the RoH ceiling of 0.3, anti-neurodiscrimination, and revocability—must be translated into formal ALN rules and Rust assertions. For example, a rule could be created that prevents a shard from being committed to the ledger without a valid, non-expired consent envelope linked to a DID [106].

The second phase involves the creation of a pilot "Phoenix" node, which will serve as the ground truth against which all proposed evolution paths are tested. Based on the provided context, a Phoenix-class MAR (Managed Aquifer Recharge) cell pilot is an ideal candidate for this role . This involves selecting a specific geographic location and defining its unique corridor parameters, such as hydraulic loading rates, PFAS concentration limits, and temperature bands, which can be encoded as a hex geoproof for immutability . This MAR cell becomes a "mirror" for the entire system; any proposed evolution path, whether it is a new planning tool, a cognitive enhancement for an engineer, or a redesign of the cell itself, must be scored on how it tightens the cell's K/E/R metrics and corridor definitions . This anchors the abstract policy in a tangible, real-world ecological problem, providing a clear and compelling metric for success. The repository layout and accompanying Rust code provided in the context serve as a strong starting point for this phase, offering a production-ready spine that can be extended and customized .

The final phase of implementation is centered on the design of the EVOLVE token mechanics, framed as a workflow metric rather than a status symbol to reinforce the anti-coercion principle . While the immediate priority is to lock in the principles, placeholder mechanics should be designed. This includes defining daily EVOLVE caps, which could be tied to an individual's safety profile and verified eco-impact contributions, and establishing per-session limits to prevent large, abrupt changes in capability . These mechanics would be managed within an `evolution.budget.v1` shard, which would be consumed by the neurorights firewall during the approval process. As the system matures and gathers data from the pilot, these parameters can be tuned iteratively. The analysis of logged evolution steps, complete with their K/E/R scores and outcomes, will provide the empirical basis for refining the caps, connector designs, and overall policy to better balance safety and progress . This phased approach, moving from abstract specification to concrete implementation, ensures that the policy is built on a solid foundation of technical and procedural rigor, ready to be deployed in a controlled environment before wider application.

Reference

1. (PDF) The DAta Protection REgulation COmpliance Model https://www.researchgate.net/publication/335953715_The_DAta_Protection_REgulation_COmpliance_Model
2. Spam-Ham NLP Basic Project https://www.kaggle.com/code/syamalakumar/spam-ham-nlp-basic-project

3. A generic non-invasive neuromotor interface for human- ... https://www.nature.com/articles/s41586-025-09255-w
4. A Survey of LLM-Driven AI Agent Communication https://arxiv.org/html/2506.19676v4
5. Arxiv今日论文| 2026-01-06 http://lonepatient.top/2026/01/06/arxiv_papers_2026-01-06.html
6. (PDF) A generic non-invasive neuromotor interface for ... https://www.researchgate.net/publication/393955605_A_generic_non-invasive_neuromotor_interface_for_human-computer_interaction
7. From passive response to active interaction and closed ... https://www.sciencedirect.com/science/article/pii/S2452199X25005031
8. Machine Learning et Interfaces 2D/3D pour automatiser ... https://enac.hal.science/tel-05337253v1/file/138047_GROSSETETE_2025.pdf
9. Ontology-Driven Multi-Agent System for Cross-Domain Art ... https://www.mdpi.com/1999-5903/17/11/517
10. Synthetic consciousness architecture https://www.frontiersin.org/journals/robotics-and-ai/articles/10.3389/frobt.2024.1437496/full
11. A2Seek: Towards Reasoning-Centric Benchmark for Aerial ... https://openreview.net/pdf/0e221ac06851beb7168298c77917846a93242ebf.pdf
12. Towards a Governance Framework for Brain Data | Neuroethics https://link.springer.com/article/10.1007/s12152-022-09498-8
13. (PDF) Neurorights, Mental Privacy, and Mind Reading https://www.researchgate.net/publication/382079309_Neurorights_Mental_Privacy_and_Mind_Reading
14. Promoting and Advancing Human Rights in Global AI ... https://www.researchgate.net/publication/389164731_Promoting_and_Advancing_Human_Rights_in_Global_AI_Ecosystems_The_Need_for_A_Comprehensive_Framework_under_International_Law
15. Common human rights challenges raised by different ... https://www.researchgate.net/publication/356412675_COMMON_HUMAN_RIGHTS_CHALLENGES_RAISED_BY_DIFFERENT_APPLICATIONS_OF_NEUROTECHNOLOGIES_IN_THE_BIOMEDICAL_FIELD
16. 333333 23135851162 the 13151942776 of 12997637966 ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/words-333333.txt
17. Hw3 Stats Google 1gram | PDF | Internet Forum | Software https://www.scribd.com/document/917428277/Hw3-Stats-Google-1gram
18. The protection of mental privacy in the area of neuroscience https://www.europarl.europa.eu/RegData/etudes/STUD/2024/757807/EPRS_STU(2024)757807_EN.pdf

19. Full article: Neurotechnologies and human rights https://www.tandfonline.com/doi/full/10.1080/13642987.2024.2310830
20. Neurotechnology https://www.ohchr.org/sites/default/files/documents/hrbodies/hrcouncil/advisorycommittee/neurotechnology/03-ngos/ac-submission-cso-neurorightsfoundation.pdf
21. Connecting the dots in trustworthy Artificial Intelligence ... https://www.sciencedirect.com/science/article/pii/S1566253523002129
22. A/HRC/57/61 - General Assembly - the United Nations https://docs.un.org/en/A/HRC/57/61
23. International Human Rights Protection Gaps in the Age of ... https://ntc.columbia.edu/wp-content/uploads/2022/05/NeurorightsFoundationPUBLICAnalysis5.6.22.pdf
24. Safety and Privacy in Immersive Extended Reality https://link.springer.com/article/10.1007/s44206-024-00114-1
25. Report of the International Bioethics Committee ... https://unesdoc.unesco.org/ark:/48223/pf0000378724
26. Antimicrobial, Oxidant, Cytotoxic, and Eco-Safety ... https://www.mdpi.com/1424-8247/18/7/976
27. Is purchasing of vegetable dishes affected by organic or ... https://www.sciencedirect.com/science/article/abs/pii/S0195666322000861
28. Hybridize Machine Learning Methods and Optimization ... https://ieeexplore.ieee.org/iel8/6287639/10820123/10851261.pdf
29. Biological Safety and Sustainability of Metal Nanomaterials ... https://www.mdpi.com/2072-666X/16/1/15
30. Hybridise Machine Learning methods and optimisation ... https://ieeexplore.ieee.org/iel8/6287639/6514899/10851261.pdf
31. Integrating RNA Interference and Nanotechnology https://www.mdpi.com/2223-7747/14/6/977
32. Mind the gap: bridging ethical considerations and regulatory ... https://pmc.ncbi.nlm.nih.gov/articles/PMC12325254/
33. Non-Invasive Brain-Computer Interfaces: State of the Art ... https://www.researchgate.net/publication/383429753_Non-invasive_Brain-Computer_Interfaces_State_of_the_Art_and_Trends
34. Ethical Considerations in Ending Exploratory Brain–Computer ... https://pmc.ncbi.nlm.nih.gov/articles/PMC8219386/
35. Validation of a Brain-Computer Interface (BCI) System ... https://www.researchgate.net/publication/318133336_Validation_of_a_Brain-

Computer_Interface_BCI_System_Designed_for_Patients_with_Disorders_of_Consciousness_DOC_Regular_and_Sham_Testing_with_Healthy_Participants

36. Words | PDF | Science | Engineering https://www.scribd.com/doc/88199361/Words
37. (PDF) The Lyapunov Neural Network: Adaptive Stability ... https://www.researchgate.net/publication/326799338_The_Lyapunov_Neural_Network_Adaptive_Stability_Certification_for_Safe_Learning_of_Dynamical_Systems
38. adaptive control of ill-defined systems - Springer Link https://link.springer.com/content/pdf/10.1007/978-1-4684-8941-5.pdf
39. Artificial Intelligence and Sustainable Energy Systems https://www.mdpi.com/topics/Artificial_Intelligence_Energy_Systems
40. (PDF) Unsupervised adaptation of an ECoG based brain– ... https://www.researchgate.net/publication/366142736_Unsupervised_adaptation_of_an_ECoG_based_brain-computer_interface_using_neural_correlates_of_task_performance
41. Applications of Intelligent Control to Engineering Systems https://www.academia.edu/41219562/Applications_of_Intelligent_Control_to_Engineering_Systems
42. Spatiotemporal characteristics of ecological resilience and ... https://www.nature.com/articles/s41598-024-67628-z
43. Critical review of existing built environment resilience ... https://www.sciencedirect.com/science/article/pii/S2212420917302388
44. Ecological resilience, robustness and vulnerability: How do ... https://www.researchgate.net/publication/259525010_Ecological_resilience_robustness_and_vulnerability_How_do_these_concepts_benefit_ecosystem_management
45. A Critical Review of Social Resilience Properties and ... https://journal.hep.com.cn/ijodrs/EN/10.1007/s13753-021-00378-y
46. The Spatiotemporal Characteristics and Spatial Linkages ... https://www.mdpi.com/2079-8954/13/9/776
47. “Resilience-based framework for enhancing NaTech risk ... https://link.springer.com/article/10.1007/s10669-025-10056-9
48. cmnt_vocab.txt https://www.cs.cmu.edu/~ark/blog-data/data/blog_data_v1_0/dk/hbc_data/data/cmnt_vocab.txt
49. Spam-ham analysis using Naive Bayes https://www.kaggle.com/code/sadiqueabdullah/spam-ham-analysis-using-naive-bayes
50. Ics - Uci.edu DMDB Chandra Enron2.1 Words | PDF https://www.scribd.com/document/675329216/Ics-uci-Edu-Dmdb-Chandra-Enron2-1-Words

51. Regulatory Overview for Neurological Devices https://www.fda.gov/medical-devices/neurological-devices/regulatory-overview-neurological-devices
52. Neurorights in the Constitution: from neurotechnology to ethics ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11491849/
53. (PDF) Neurotechnologies in the AI Act: Moving away from ... https://www.researchgate.net/publication/399212126_Neurotechnologies_in_the_AI_Act_Moving_away_from_the_Neurorights_Debate
54. Neurorights (Chapter 26) - The Cambridge Handbook of ... https://www.cambridge.org/core/books/cambridge-handbook-of-the-right-to-freedom-of-thought/neurorights/B1AEF25AD18D9C8164CE9B366979B664
55. Chilean Supreme Court ruling on the protection of brain ... https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2024.1330439/full
56. A framework to explore the effects of urban planning ... https://www.sciencedirect.com/science/article/pii/S2212041618300032
57. Full article: Regenerative sustainability. A relational model ... https://www.tandfonline.com/doi/full/10.1080/23251042.2023.2239538
58. An Ecological Dynamics Approach to Understanding ... https://www.mdpi.com/1660-4601/19/6/3691
59. AI for brain-computer interfaces | Request PDF https://www.researchgate.net/publication/379676293_AI_for_brain-computer_interfaces
60. Download book PDF https://link.springer.com/content/pdf/10.1007/978-3-642-21344-1.pdf
61. Review Large language models for battery prognostics https://www.sciencedirect.com/science/article/pii/S2095495625009660
62. Indicators and methods for the ecological status ... https://www.researchgate.net/profile/Laurence-Carvalho/publication/279642003_Nutrients_and_eutrophication_in_lakes/links/559d34bc08ae04e36508bcf2/Nutrients-and-eutrophication-in-lakes.pdf
63. Using problem formulation for fit‐for‐purpose pre‐market ... https://efsa.onlinelibrary.wiley.com/doi/10.2903/j.efsa.2019.e170708
64. Regulating neural data processing in the age of BCIs https://pmc.ncbi.nlm.nih.gov/articles/PMC11951885/
65. “You shall have the thought”: habeas cogitationem as a New ... https://link.springer.com/article/10.1007/s12152-024-09551-8
66. ISO/IEC JTC 1/SC 42 - Artificial intelligence https://www.iso.org/committee/6794475.html

67. Developments and Advances in Defense and Security https://link.springer.com/content/pdf/10.1007/978-3-319-78605-6.pdf
68. Liste Von Abkürzungen BR | PDF | Telecommunications https://www.scribd.com/document/841150772/Liste-Von-Abkurzungen-Br
69. Advances in Wireless Networks and Information Systems https://link.springer.com/content/pdf/10.1007/978-3-642-14350-2.pdf
70. Artificial intelligence - Standards by ISO/IEC JTC 1/SC 42 https://www.iso.org/committee/6794475/x/catalogue/p/%200/u/1/w/0/d/0
71. Advances On Digital Television and Wireless Multimedia ... https://www.scribd.com/document/451572091/Advances-on-Digital-Television-and-Wireless-Multimedia-Communications-pdf
72. Arxiv今日论文| 2026-01-19 http://lonepatient.top/2026/01/19/arxiv_papers_2026-01-19.html
73. 人工智能2025_7_15 http://arxivdaily.com/thread/69431
74. EN Horizon Europe Work Programme 2025 8. Climate, Energy ... https://research-and-innovation.ec.europa.eu/document/download/64e1031e-7814-4c11-ab14-a22b1f5236c2_en
75. Analysis and Modelling of Severe Accidents for Liquid Metal ... https://www-pub.iaea.org/MTCD/publications/PDF/TE-2079web.pdf
76. Current and Future Approaches to AI Governance (Part II) https://www.cambridge.org/core/books/cambridge-handbook-of-responsible-artificial-intelligence/current-and-future-approaches-to-ai-governance/4F9A8DEE6D49F891ADE37F8FA1A9BD49
77. Technologies and Applications for Big Data Value https://oadoi.org/10.1007/978-3-030-78307-5
78. ISO/IEC JTC 1/SC 42/WG 4 N 238 https://www.itu.int/en/ITU-T/focusgroups/ai4h/Documents/all/FGAI4H-H-025-A02.pdf
79. ISO/IEC JTC 1/SC 42 - Artificial intelligence https://www.iso.org/cms/%20render/live/es/sites/isoorg/contents/data/committee/67/94/6794475/x/catalogue/
80. Chilean neurorights legislation and its relevance for mental ... https://www.researchgate.net/publication/373903709_Chilean_neurorights_legislation_and_its_relevance_for_mental_health_Criticisms_and_outlook
81. Chilean Supreme Court ruling on the protection of brain ... https://www.academia.edu/115498900/Chilean_Supreme_Court_ruling_on_the_protection_of_brain_activity_neurorights_personal_data_protection_and_neurodata

82. Neurotechnology and Children's Rights — Preparing for ... https://www.unicef.org/innocenti/media/11426/file/UNICEF-Innocenti-Neurotech-Childrens-rights-Report-2025.pdf
83. Data analytics subgroup report - EMA https://www.ema.europa.eu/en/documents/report/data-analytics-subgroup-report_en.pdf
84. (PDF) Brain-MCP: A Fully Homomorphic Neuro-Symbolic ... https://www.researchgate.net/publication/393129669_Brain-MCP_A_Fully_Homomorphic_Neuro-Symbolic_Protocol_for_Governing_System_3_Cognition_and_Restoring_Semantic_Scarcity
85. Open Semantic Technologies for Intelligent System https://link.springer.com/content/pdf/10.1007%2F978-3-030-60447-9.pdf
86. K090957 - 510(k) Premarket Notification - FDA https://www.accessdata.fda.gov/scripts/cdrh/cfdocs/cfpmn/pmn.cfm?ID=K090957
87. Premarket Approval (PMA) - FDA https://www.accessdata.fda.gov/scripts/cdrh/cfdocs/cfpma/pma.cfm?id=P960009S229
88. Premarket Approval (PMA) - FDA https://www.accessdata.fda.gov/scripts/cdrh/cfdocs/cfpma/pma.cfm?id=p960009s219
89. Premarket Approval (PMA) - FDA https://www.accessdata.fda.gov/scripts/cdrh/cfdocs/cfpma/pma.cfm?id=P960009S478
90. NeuroPort Electrode 96, Electrodes | Products https://blackrockneurotech.com/products/neuroport-electrode/
91. P960009/S478 Summary of Safety and Effectiveness (SSED) https://www.accessdata.fda.gov/cdrh_docs/pdf/P960009S478B.pdf
92. Premarket Approval (PMA) - FDA https://www.accessdata.fda.gov/scripts/cdrh/cfdocs/cfpma/pma.cfm?ID=P960009S036
93. 510(k) Premarket Notification - accessdata.fda.gov https://www.accessdata.fda.gov/scripts/cdrh/cfdocs/cfPMN/pmn.cfm?ID=4629
94. K202174.pdf - accessdata.fda.gov https://www.accessdata.fda.gov/cdrh_docs/pdf20/K202174.pdf
95. A Critical Perspective on NeuroRights https://pmc.ncbi.nlm.nih.gov/articles/PMC8573066/
96. A Threat to Autonomy? The Intrusion of Predictive Brain ... https://pmc.ncbi.nlm.nih.gov/articles/PMC4685604/
97. The Land of Speculative Ethics and Alarming Claims? https://www.researchgate.net/publication/379537529_Neurorights_The_Land_of_Speculative_Ethics_and_Alarming_Claims

98. Bayesian Reanalysis of a Randomized Controlled Trial https://pubmed.ncbi.nlm.nih.gov/37684058/
99. Novel Neurorights: From Nonsense to Substance - PMC https://pmc.ncbi.nlm.nih.gov/articles/PMC8821782/
100. A Bayesian non-inferiority approach using experts' margin ... https://pmc.ncbi.nlm.nih.gov/articles/PMC6751616/
101. Robotic Assistance for Posterior Approach Total Hip ... https://pubmed.ncbi.nlm.nih.gov/35124193/
102. PMF IAS Polity CA Compilation for UPSC CSE Prelims 2024 https://pmfstore.s3.ap-south-1.amazonaws.com/PMFIAS-CA-Prelims-2024-PF-Polity.pdf?X-Amz-Content-Sha256=UNSIGNED-PAYLOAD&X-Amz-Algorithm=AWS4-HMAC-SHA256&X-Amz-Credential=AKIAIWGL56NSK6RLSKQA%2F20251212%2Fap-south-1%2Fs3%2Faws4_request&X-Amz-Date=20251212T190128Z&X-Amz-SignedHeaders=host&X-Amz-Expires=72000&X-Amz-Signature=1cb81ab9aff8078a712825ddb40132699093ff1728a77c1de08884e69622f82f
103. GLOTEL | PDF | Modulation | Antenne (Radio) https://fr.scribd.com/document/506194637/GLOTEL
104. Experts for Eg-Digital-Library-Final-Wrapper https://www.linknovate.com/search/?query=eg-digital-library-final-wrapper
105. Apple Platform Security https://help.apple.com/pdf/security/en_US/apple-platform-security-guide.pdf
106. A Novel Zero-Trust Identity Framework for Agentic AI https://arxiv.org/html/2505.19301v2